

# DS2 Final Project

Wanxin Qi, Lesi He, Ke Xu

```
library(tidyverse)
library(lubridate)
library(caret)
library(ggplot2)
library(corrplot)
library(vip)
library(rpart.plot)
library(ranger)
library(GGally)
library(pdp)

knitr::opts_chunk$set(echo = TRUE, message = FALSE, warning = FALSE)

theme_set(theme_minimal() + theme(legend.position = "bottom"))
options(
  ggplot2.continuous.colour = "viridis",
  ggplot2.continuous.fill = "viridis")
scale_colour_discrete = scale_colour_viridis_d
scale_fill_discrete = scale_fill_viridis_d
```

## Data

```
bike = read.csv("./data/SeoulBikeData.csv", check.names = F)

# Missing value in the dataset
sum(is.na(bike))
```

```
## [1] 0
```

```
# All the 0 hourly rented bike count are non-functioning day
bike %>%
  janitor::clean_names() %>%
  filter(rented_bike_count == 0) %>%
  count(functioning_day)
```

```
##   functioning_day    n
## 1                No 295
```

```
# tidy
bike = bike %>%
```

```

janitor::clean_names() %>%
mutate(
  date = dmy(date),
  week = weekdays(date, abbreviate = TRUE),
  week = factor(week, levels = c("Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun")),
  rent = as.numeric(rented_bike_count),
  hour = as.numeric(hour),
  temp = as.numeric(temperature_c),
  hum = as.numeric(humidity_percent),
  wind = as.numeric(wind_speed_m_s),
  visibility = as.numeric(visibility_10m),
  dew_temp = as.numeric(dew_point_temperature_c),
  radiation = as.numeric(solar_radiation_mj_m2),
  rain = as.numeric(rainfall_mm),
  snow = as.numeric(snowfall_cm),
  season = as.factor(seasons),
  holiday = as.factor(ifelse(holiday == "No Holiday", "No", "Yes")),
  func = as.factor(functioning_day)
) %>%
select(rent, hour, temp, hum, wind, visibility, dew_temp, radiation,
       rain, snow, season, week, holiday, func)

# Dataset of the research
set.seed(2022)
bike = bike[sample(nrow(bike), 1000),]

# Partition
set.seed(2)
trainRows = createDataPartition(y = bike$rent, p = 0.8, list = FALSE)
trainData = bike[trainRows,]
testData = bike[-trainRows,]

train_x = model.matrix(rent ~ ., bike)[trainRows, -1]
train_y = bike$rent[trainRows]
test_x = model.matrix(rent ~ ., bike)[-trainRows, -1]
test_y = bike$rent[-trainRows]

```

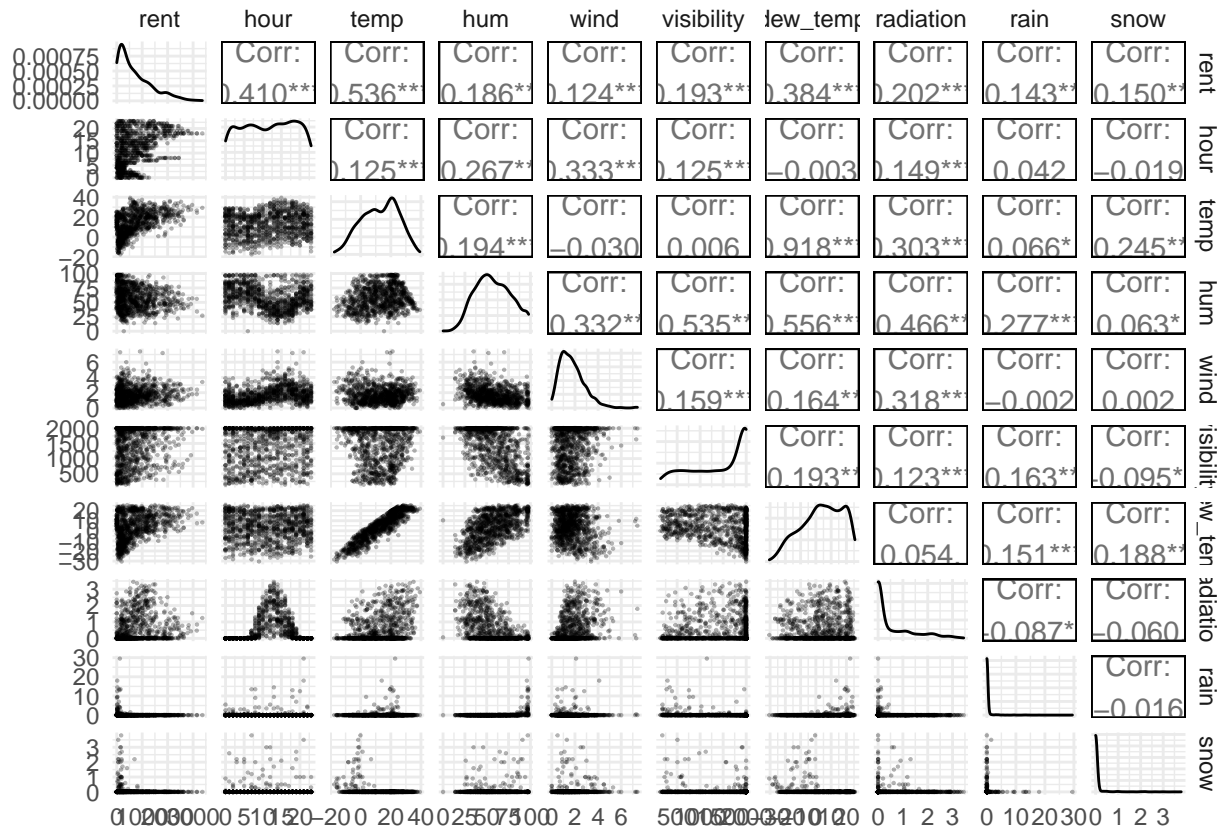
## EDA

```

# Summary of Data
# knitr::kable(summary(bike), digits = 2)

# Correlation and Sactter Plot
cont = bike %>% select(-week, -holiday, -func, -season)
ggpairs(cont, lower = list(continuous = wrap("points", alpha = 0.3, size = 0.1)))

```



# Average Hourly Rental Bike Count Across Seasons

bike %>%

mutate(hour = as.factor(hour)) %>%

group\_by(season, hour) %>%

summarise(rent.avg = mean(rent)) %>%

mutate(hour = as.integer(hour)) %>%

ggplot(aes(x = hour, y = rent.avg)) +

geom\_point(aes(color = season)) +

geom\_line(aes(color = season)) +


labs(

title = "Average Hourly Rented Bike Count Across Seasons",

x = "Hour",

y = "Average Hourly Rented Bike Count")

Average Hourly Rented Bike Count Across Seasons



```
# Average Hourly Rental Bike Count Across Weekdays
```

```
bike %>%
```

```
  mutate(hour = as.factor(hour)) %>%
```

```
  group_by(week, hour) %>%
```

```
  summarise(rent.avg = mean(rent)) %>%
```

```
  mutate(hour = as.numeric(hour)) %>%
```

```
  ggplot(aes(x = hour, y = rent.avg)) +
```

```
  geom_point(aes(color = week)) +
```

```
  geom_line(aes(color = week), alpha = 0.5) +
```


```
  labs(
```

```
    title = "Average Hourly Rented Bike Count Across Weekdays",
```

```
    x = "Hour",
```

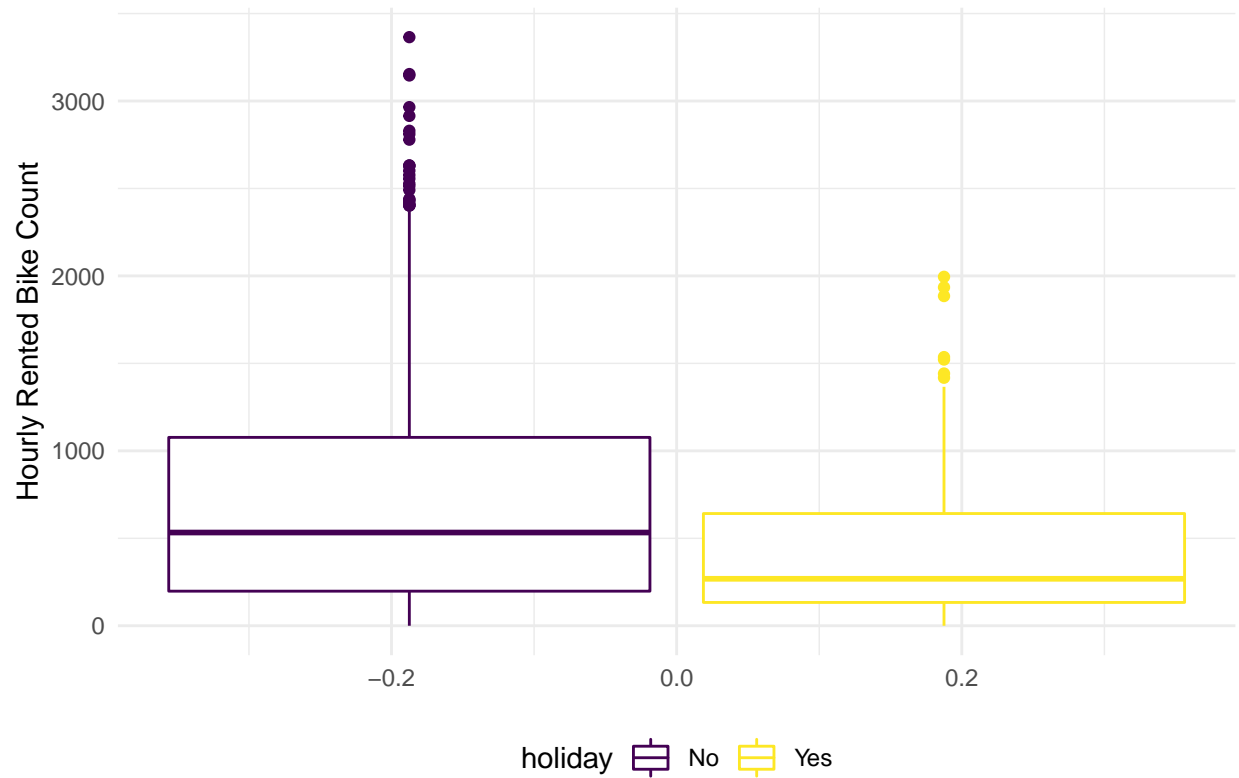
```
    y = "Average Hourly Rented Bike Count")
```

Average Hourly Rented Bike Count Across Weekdays



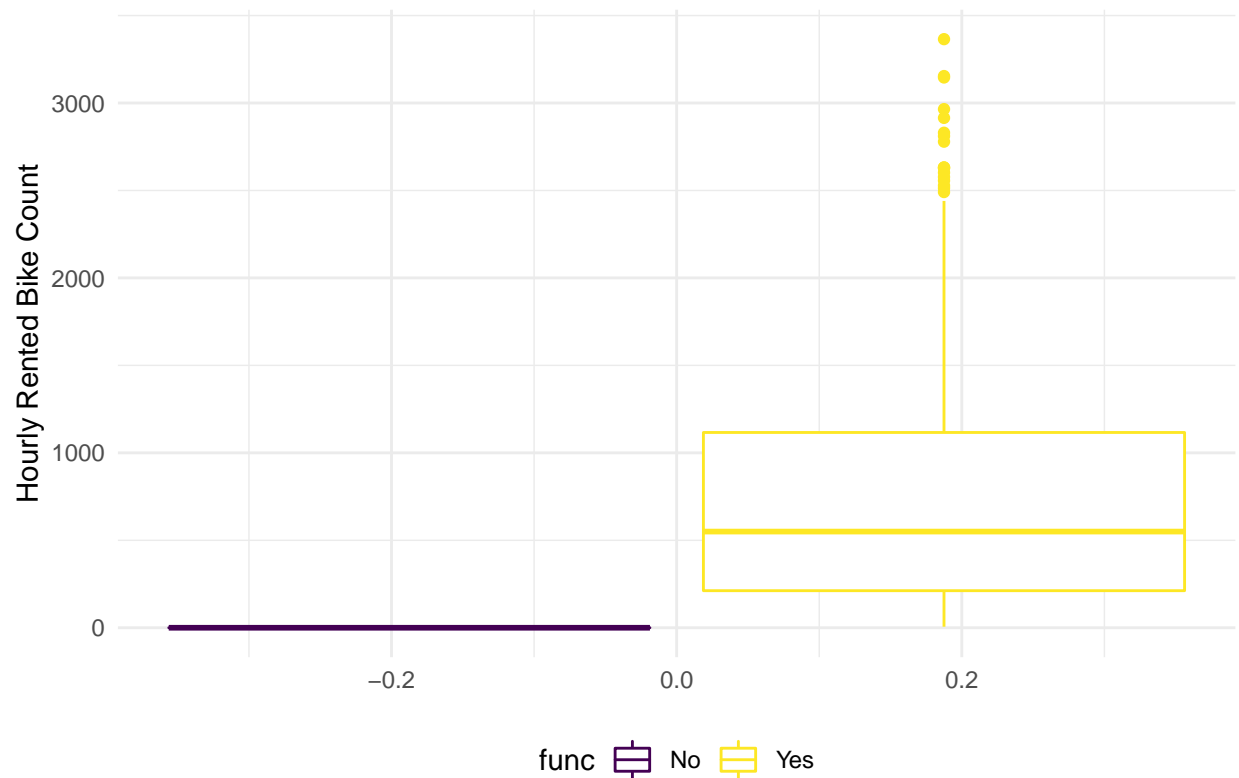
```
# Hourly Rented Bike Count Across Holiday
bike %>%
  ggplot(aes(y = rent, color = holiday)) +
  geom_boxplot() +
  labs(
    title = "Hourly Rented Bike Count Across Holiday",
    y = "Hourly Rented Bike Count")
```

Hourly Rented Bike Count Across Holiday



```
# Hourly Rented Bike Count Across Functional Day
bike %>%
  ggplot(aes(y = rent, color = func)) +
  geom_boxplot() +
  labs(
    title = "Hourly Rented Bike Count Across Functional Day",
    y = "Hourly Rented Bike Count")
```

## Hourly Rented Bike Count Across Functional Day



## Modeling

LM

```
# Resampling Method - 10-Fold CV
ctrl1 = trainControl(method = "cv", number = 10)

set.seed(2)
lm.fit = train(train_x, train_y,
               method = "lm",
               trControl = ctrl1)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1120.55  -270.93   -55.25   215.50  1724.37
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2.461e+02  3.476e+02  -0.708  0.479248
## hour        2.748e+01  2.559e+00  10.739 < 2e-16 ***
```

```
## temp      2.630e+01  1.292e+01   2.036 0.042092 *
## hum       -8.828e+00  3.661e+00  -2.411 0.016121 *
## wind      1.809e+01  1.723e+01   1.050 0.294062
## visibility -9.858e-03  3.365e-02  -0.293 0.769627
## dew_temp   2.371e+00  1.360e+01   0.174 0.861649
## radiation  -1.220e+02  2.666e+01  -4.575 5.54e-06 ***
## rain      -6.039e+01  9.962e+00  -6.062 2.09e-09 ***
## snow      5.220e+01  4.899e+01   1.066 0.286949
## seasonSpring -2.230e+02  4.792e+01  -4.653 3.84e-06 ***
## seasonSummer -1.768e+02  5.837e+01  -3.029 0.002535 **
## seasonWinter -3.780e+02  6.477e+01  -5.837 7.79e-09 ***
## weekTue    1.112e+02  5.567e+01   1.998 0.046043 *
## weekWed    9.210e+00  5.601e+01   0.164 0.869416
## weekThu    2.041e+02  5.940e+01   3.436 0.000622 ***
## weekFri    1.018e+02  5.751e+01   1.770 0.077156 .
## weekSat   -1.983e+01  5.965e+01  -0.332 0.739611
## weekSun   -1.099e+02  5.898e+01  -1.863 0.062822 .
## holidayYes -1.089e+02  7.370e+01  -1.478 0.139736
## funcYes    1.057e+03  1.034e+02  10.225 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 436.4 on 779 degrees of freedom
## Multiple R-squared:  0.5747, Adjusted R-squared:  0.5638
## F-statistic: 52.64 on 20 and 779 DF,  p-value: < 2.2e-16
```

```
# test error
```

```
lm.pred = predict(lm.fit, newdata = test_x)
RMSE(lm.pred, test_y)
```

```
## [1] 396.6971
```


```
# 396.6971
```

LASSO

```
set.seed(2)
lasso.fit = train(train_x, train_y,
                  method = "glmnet",
                  tuneGrid = expand.grid(alpha = 1,
                                         lambda = exp(seq(5, 0, length = 100))),
                  trControl = ctrl1)
plot(lasso.fit, xTrans = log, main = "Tuning Process of LASSO")
```



## Tuning Process of LASSO



```
lasso.fit$bestTune
```

```
##      alpha  lambda  
## 18         1 2.359821
```

```
# lambda = 2.359821
```

```
coef(lasso.fit$finalModel, lasso.fit$bestTune$lambda)
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s1  
## (Intercept) -304.00576  
## hour        27.71575  
## temp        27.33860  
## hum         -7.74392  
## wind        13.95456  
## visibility   .  
## dew_temp     .  
## radiation   -110.63904  
## rain        -59.52354  
## snow        38.81008  
## seasonSpring -203.01668  
## seasonSummer -147.45155  
## seasonWinter -367.56817  
## weekTue      96.17458
```

```
## weekWed      .
## weekThu      189.14731
## weekFri       86.33245
## weekSat      -24.14187
## weekSun     -113.71925
## holidayYes   -95.42776
## funcYes     1025.85274
```

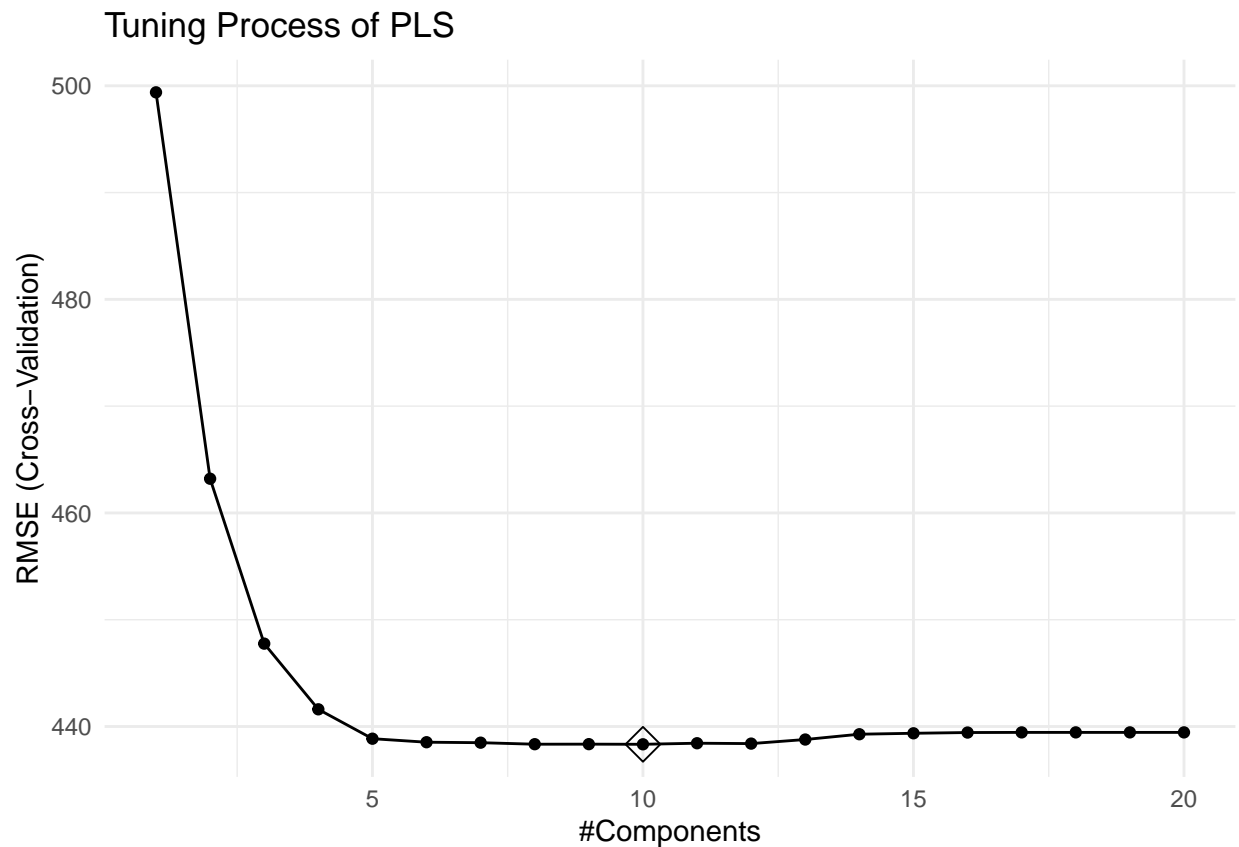
```
# test error
lasso.pred = predict(lasso.fit, newdata = test_x)
RMSE(lasso.pred, test_y)
```

```
## [1] 396.3537
```

```
# 396.3537
```

PLS

```
set.seed(2)
pls.fit = train(train_x, train_y, method = "pls",
                tuneGrid = data.frame(ncomp = 1:20),
                trControl = ctrl1,
                preProcess = c("center", "scale"))
ggplot(pls.fit, highlight = TRUE) +
  labs(title = "Tuning Process of PLS")
```



```
pls.fit$bestTune
```

```
##      ncomp  
## 10      10
```

```
# ncomp = 10
```

```
summary(pls.fit$finalModel)
```

```
## Data:      X dimension: 800 20  
## Y dimension: 800 1  
## Fit method: oscorespls  
## Number of components considered: 10  
## TRAINING: % variance explained  
##           1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  
## X           14.34   25.24   32.30   36.07   39.59   43.87   47.86  
## .outcome    43.44   51.91   55.62   56.82   57.35   57.42   57.45  
##           8 comps  9 comps 10 comps  
## X           52.00   56.00   60.14  
## .outcome    57.45   57.46   57.46
```

```
# 57.46% of variance explained
```

```
# test error
```

```
pls.pred = predict(pls.fit, newdata = test_x, ncomp = ncomp.cv)  
RMSE(pls.pred, test_y)
```

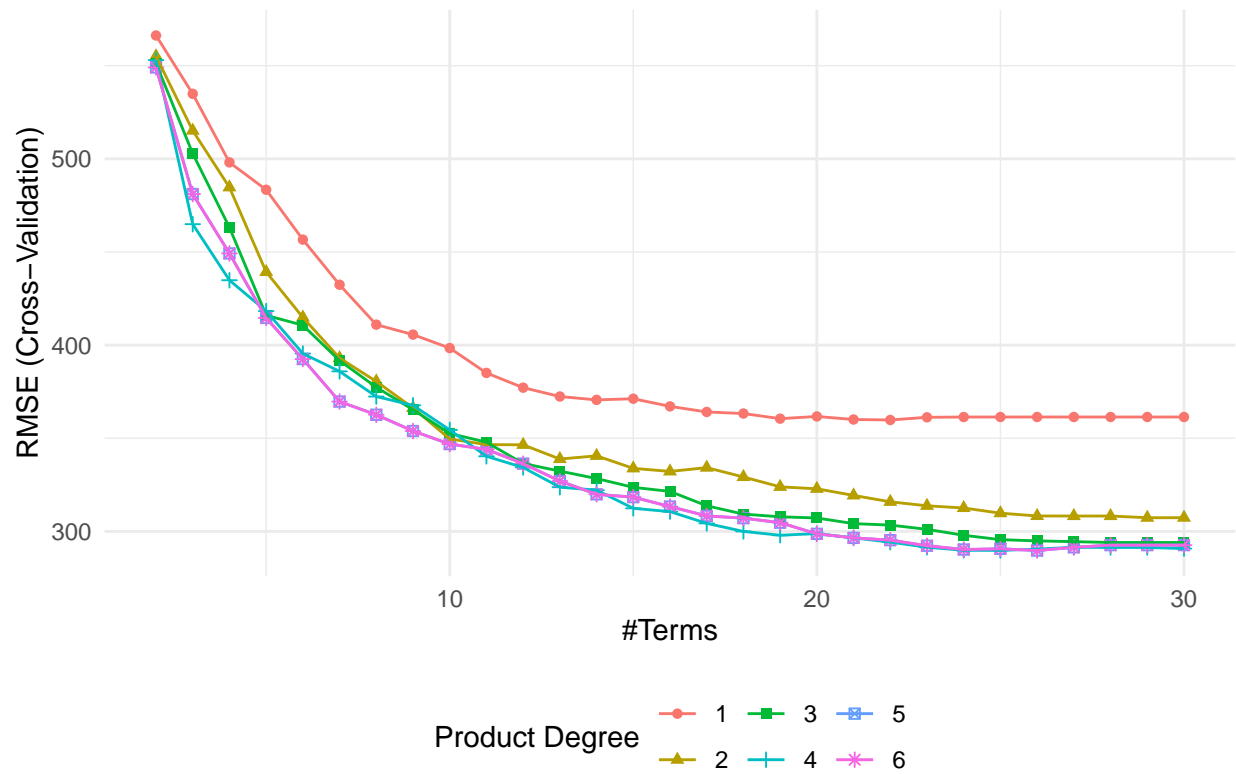
```
## [1] 395.9891
```

```
# 395.9891
```

MARS

```
set.seed(2)  
mars_grid = expand.grid(degree = 1:6, nprune = 2:30)  
mars.fit = train(train_x, train_y, method = "earth",  
                 tuneGrid = mars_grid,  
                 trControl = ctrl1)  
ggplot(mars.fit) +  
  labs(title = "Tuning Process of MARS")
```

## Tuning Process of MARS



```
mars.fit$bestTune
```

```
##      nprune degree
## 141      26      5
```

```
# nprune = 26; degree = 5
```

```
coef(mars.fit$finalModel)
```

```
##              (Intercept)
##              -6.555446e+03
##              h(33.9-temp)
##              -5.868188e+01
##              h(hour-21)
##              -4.923851e+02
##              h(21-hour)
##              5.277675e+01
##              h(21-hour) * h(28.8-temp)
##              5.696055e+00
##              h(rain-1.5)
##              -4.671311e+03
##              h(1.5-rain)
##              4.185633e+03
##              h(1.5-rain) * funcYes
##              3.210956e+02
```

```


##          h(hour-4) * h(1.5-rain) * funcYes
##          6.536085e+01
##          h(4-hour) * h(1.5-rain) * funcYes
##          9.545411e+01
##          h(8-hour) * h(33.9-temp)
##          -8.846630e+00
##          h(hour-4) * h(temp-26.9) * h(1.5-rain) * funcYes
##          -6.793200e+00
##          h(hour-4) * h(26.9-temp) * h(1.5-rain) * funcYes
##          -2.482755e+00
##          h(hour-8) * h(33.9-temp) * h(9.5-dew_temp)
##          6.892964e-02
##          h(hour-14)
##          1.579968e+02
##          h(radiation-0.03)
##          -6.203958e+03
##          h(0.03-radiation)
##          -9.935974e+03
##          h(hour-9) * h(33.9-temp)
##          3.203867e+00
##          h(17-hour) * h(radiation-0.03)
##          7.738358e+02
##          h(hour-9) * h(radiation-0.03)
##          7.833004e+02
##          h(rain-0.1)
##          4.670069e+03
##          h(1.5-rain) * seasonWinter
##          -2.875902e+02
##          h(33.9-temp) * seasonSpring
##          -1.176064e+01
##          h(1.5-rain) * weekSun * funcYes
##          -1.092494e+02
##          h(hour-8) * h(33.9-temp) * h(visibility-460) * h(dew_temp-9.5)
##          -3.339103e-04
##          h(hour-8) * h(33.9-temp) * h(460-visibility) * h(dew_temp-9.5)
##          -3.234346e-03

```

```

# Variable Importance Plot
vip(mars.fit$finalModel)

```




```
# test error
mars.pred = predict(mars.fit, newdata = test_x)
RMSE(mars.pred, test_y)
```

```
## [1] 306.2679
```

```
# 306.2679
```

Regression Tree

```
set.seed(2)
rpart.fit = train(rent ~.,
                  trainData,
                  method = "rpart",
                  tuneGrid = data.frame(cp = exp(seq(-10, -5, length = 50))),
                  trControl = ctrl1)
ggplot(rpart.fit, highlight = TRUE)
```



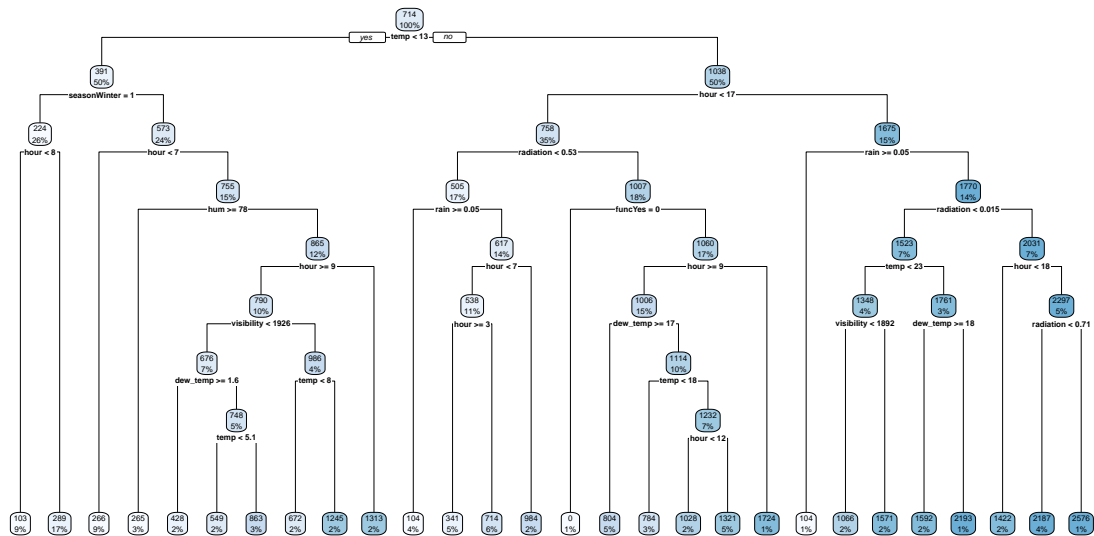
```
rpart.fit$bestTune
```

```
##          cp
## 41 0.002689588
```

```
# cp = 0.002689588
```

```
# Rpart Plot
```

```
rpart.plot(rpart.fit$finalModel)
```



```
# test error
RMSE(predict(rpart.fit, newdata = testData), test_y)
```

```
## [1] 412.142
```

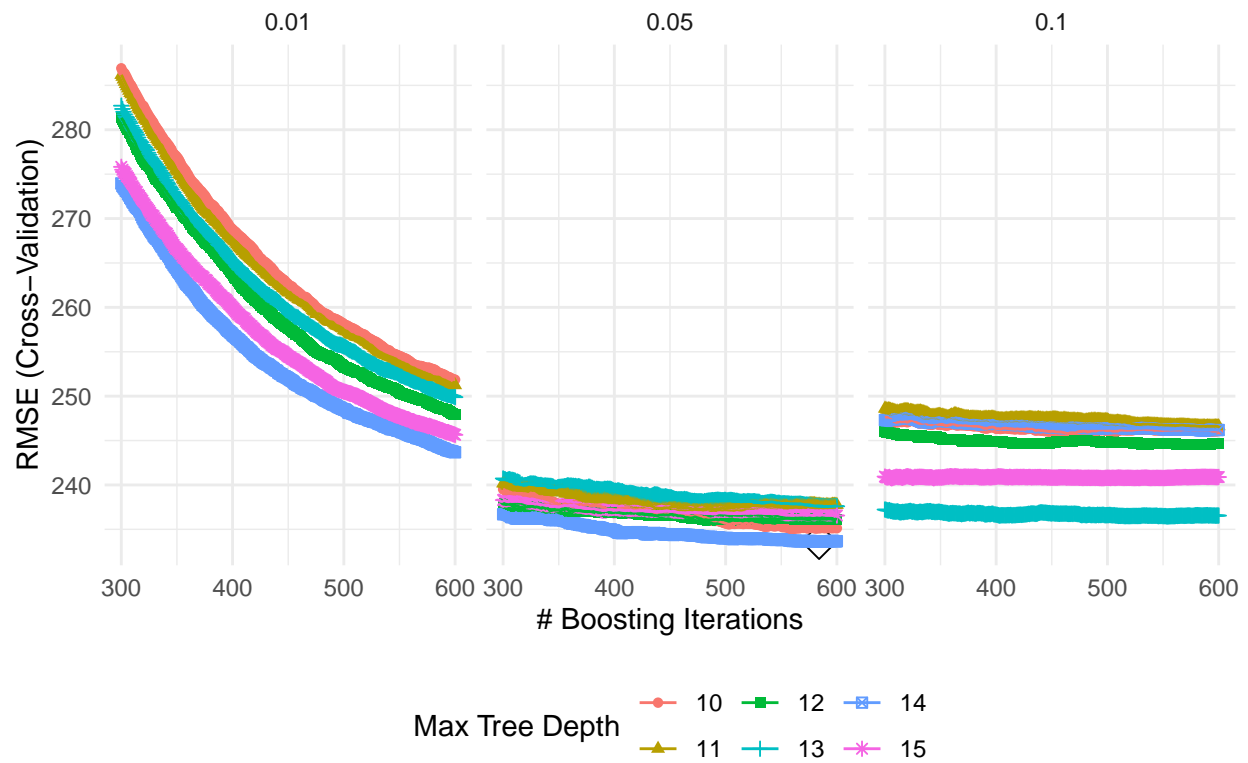
```
# 412.142
```

GBM

```
set.seed(2)
gbm.grid = expand.grid(n.trees = 300:600,
                      interaction.depth = 10:15,
                      shrinkage = c(0.01, 0.05, 0.1),
                      n.minobsinnode = 1)
gbm.fit = train(rent ~.,
               trainData,
               method = "gbm",
               tuneGrid = gbm.grid,
               trControl = ctrl1,
               verbose = FALSE)
ggplot(gbm.fit, highlight = TRUE) +
  labs(title = "Tuning Process of GBM")
```



## Tuning Process of GBM




```
gbm.fit$bestTune
```

```
##      n.trees interaction.depth shrinkage n.minobsinnode
## 3295      584              14      0.05              1
```

```
# n.trees = 584; interaction.depth = 14; shrinkage = 0.05
```

```
# Variable Importance
```

```
summary(gbm.fit$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```



```
##           var      rel.inf
## temp      temp 28.10098421
## hour      hour 27.86116128
## radiation radiation 9.72773812
## hum       hum  6.22602146
## rain      rain  5.44645569
## funcYes   funcYes 4.68396078
## dew_temp  dew_temp 3.70596063
## visibility visibility 3.39116107
## wind      wind  3.10819802
## seasonWinter seasonWinter 2.48264133
## seasonSpring seasonSpring 1.30705151
## weekSat   weekSat  0.95481430
## weekSun   weekSun  0.80497311
## weekThu   weekThu  0.73593503
## holidayYes holidayYes 0.61089943
## weekWed   weekWed  0.25528855
## weekTue   weekTue  0.23021666
## weekFri   weekFri  0.20837004
## seasonSummer seasonSummer 0.13614472
## snow      snow    0.02202402
```

```
# test error
gbm.pred = predict(gbm.fit, newdata = testData)
RMSE(gbm.pred, test_y)
```

```
## [1] 235.5603
```


```
# 235.5603
```

## Chosing Model - GBM

```
resamp = resamples(list(  
  lm = lm.fit,  
  lasso = lasso.fit,  
  pls = pls.fit,  
  mars = mars.fit,  
  gbm = gbm.fit))  
  
summary(resamp)
```

```
##  
## Call:  
## summary.resamples(object = resamp)  
##  
## Models: lm, lasso, pls, mars, gbm  
## Number of resamples: 10  
##  
## MAE  
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's  
## lm      279.8497 301.7003 345.2637 331.8926 357.1541 373.1838    0  
## lasso   282.1144 298.6512 343.1648 330.2047 353.4877 371.5051    0  
## pls     279.5372 301.2391 344.9630 331.2658 356.0821 372.9333    0  
## mars    174.1918 195.3364 201.9652 201.1253 210.5305 228.2421    0  
## gbm     127.7953 138.0474 143.3814 154.2359 173.6313 201.2062    0  
##  
## RMSE  
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's  
## lm      346.4419 382.0271 458.7146 439.4480 486.4834 524.7071    0  
## lasso   348.2079 383.3249 457.5915 438.0258 482.2165 523.4410    0  
## pls     346.3796 382.3977 457.8555 438.3322 484.7864 523.1861    0  
## mars    229.6639 261.4277 281.0772 289.7088 309.7883 370.4141    0  
## gbm     173.2382 199.9303 208.5980 233.6177 268.2053 332.9073    0  
##  
## Rsquared  
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's  
## lm      0.4356089 0.4832034 0.5507070 0.5600693 0.6341620 0.7073765    0  
## lasso   0.4387356 0.4897107 0.5538475 0.5624234 0.6327478 0.7050111    0  
## pls     0.4388923 0.4881616 0.5522216 0.5620737 0.6341998 0.7074456    0  
## mars    0.6974273 0.8068075 0.8230090 0.8089245 0.8480190 0.8707864    0  
## gbm     0.7130399 0.8616793 0.9016891 0.8693765 0.9168536 0.9292641    0
```

```
bwplot(resamp, metric = "RMSE")
```



Since the final model is GBM, which is a black-box model...

## Black-Box

Partial Dependence Plots

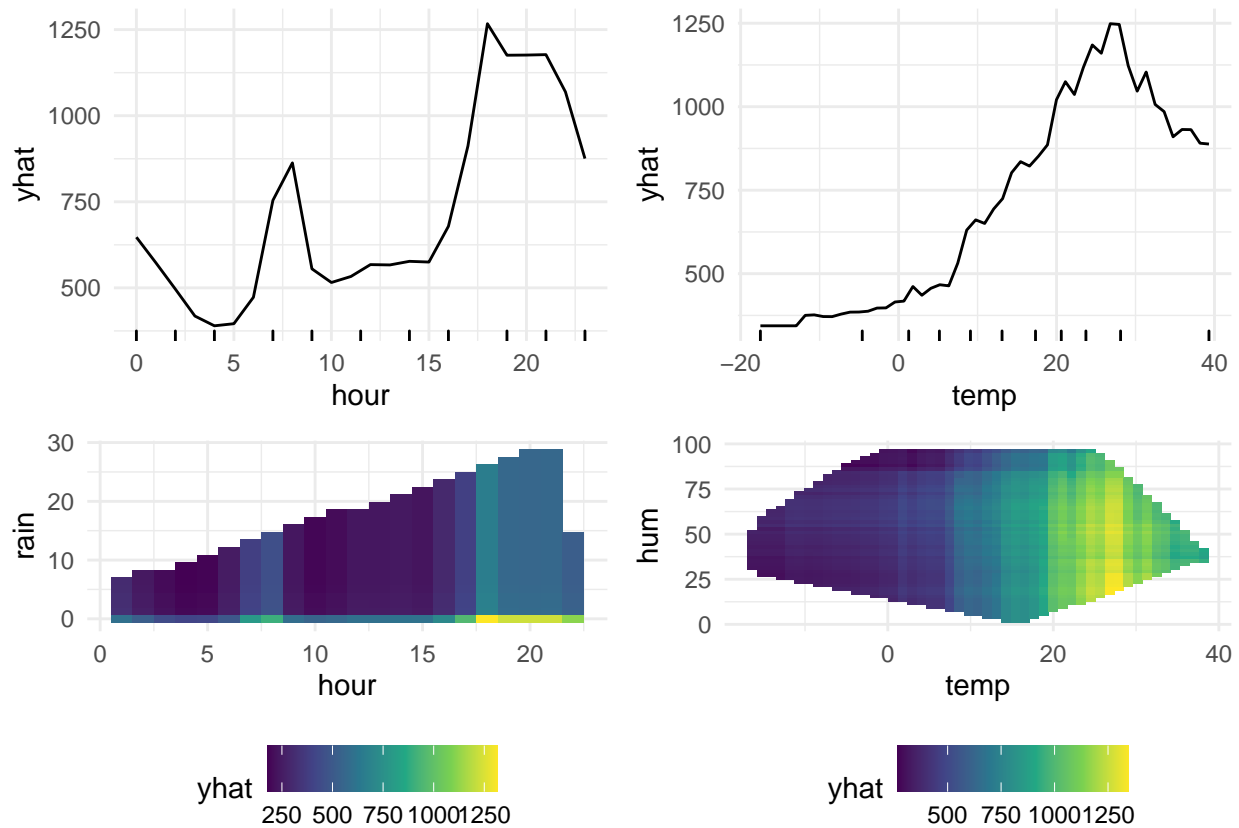
```
pdp1.gbm = gbm.fit %>%
  partial(pred.var = c("hour")) %>%
  autoplot(train = trainData, rug = TRUE)

pdp2.gbm = gbm.fit %>%
  partial(pred.var = c("temp")) %>%
  autoplot(train = trainData, rug = TRUE)

pdp3.gbm = gbm.fit %>%
  partial(pred.var = c("hour", "rain"), chull = TRUE) %>%
  autoplot(train = trainData, rug = TRUE)

pdp4.gbm = gbm.fit %>%
  partial(pred.var = c("temp", "hum"), chull = TRUE) %>%
  autoplot(train = trainData, rug = TRUE)

grid.arrange(pdp1.gbm, pdp2.gbm, pdp3.gbm, pdp4.gbm, ncol = 2, nrow = 2)
```




ICE curves

```
ice1.gbm <- gbm.fit %>%
  partial(pred.var = "hour",
    grid.resolution = 100,
    ice = TRUE) %>%
  autoplot(train = bike, alpha = .1,
    center = TRUE) +
  ggtitle("ICE - Hour, centered")

ice2.gbm <- gbm.fit %>%
  partial(pred.var = "temp",
    grid.resolution = 100,
    ice = TRUE) %>%
  autoplot(train = bike, alpha = .1,
    center = TRUE) +
  ggtitle("ICE - Temperature, centered")

grid.arrange(ice1.gbm, ice2.gbm, nrow = 1)
```

ICE – Hour, centered



ICE – Temperature, centered

